

Nate Le
CS 479
2/4/2024

Introduction

In our task, we aimed to craft 64-bit x86 Linux shellcode. It's designed to use the `execve` system function. The job is to transform the existing process into a terminal shell like `/bin/sh`. We'll dissect the assembly code, go through each line. We'll also talk about the shellcode length.

Shellcode Explanation:

```
.section .text
.globl _start
_start:
    xorl %eax, %eax      # Clear %eax register
    mov $59, %eax        # syscall number for execve (59 in 64-bit)
    lea shell_path(%rip), %rdi # address of "/bin/sh" (using RIP-relative addressing)
    xorl %esi, %esi      # null for command-line arguments
    xorl %edx, %edx      # null for environment variables
    syscall              # Invoke execve syscall

    xorl %edi, %edi      # Clear %edi register (exit status)
    mov $60, %eax        # syscall number for exit (60 in 64-bit)
    syscall              # Invoke exit syscall

.section .data
shell_path: .asciz "/bin/sh"
```

This 64-bit x86 Linux assembly shellcode, organized into sections `.text` and `.data`, begins with the `_start` label, the program's entry point. It initializes the `%eax` register to zero, representing the syscall number. Then, it loads the syscall number for `execve` (59 in 64-bit) into `%eax` and the address of the string `"/bin/sh"` into `%rdi` using RIP-relative addressing. Following this, it clears `%esi` and `%edx`, preparing them for null command-line arguments and environment variables. The `syscall` instruction invokes the `execve` syscall, executing `/bin/sh`. Next, it clears `%edi` to zero and loads the syscall number for `exit` (60 in 64-bit) into `%eax`, followed by another `syscall` to exit the program. Finally, the `.data` section defines the `shell_path` label, pointing to the string `"/bin/sh"`, essential for the `execve` syscall. This shellcode effectively spawns a shell process and then exits.

Shellcode length

I use python script to find my shellcode bytes, this is the result

My shellcode is 435 bytes long.

Here they are: -- 2E 74 65 78 74 0A 20 20 20 20 2E 67 6C 6F 62 61 6C 20 5F 73 74 61 72 74
0A 5F 73 74 61 72 74 3A 0A 20 20 20 20 78 6F 72 6C 20 25 65 61 78 2C 20 25 65 61 78 0A 20
20 20 20 6D 6F 76 62 20 24 31 31 2C 20 25 61 6C 20 20 20 20 20 20 20 20 20 20 23 20 73
79 73 63 61 6C 6C 20 6E 75 6D 62 65 72 20 66 6F 72 20 65 78 65 63 76 65 0A 20 20 20 20 6C
65 61 6C 20 73 68 65 6C 6C 5F 70 61 74 68 2C 20 25 65 62 78 20 20 20 23 20 61 64 64 72 65
73 73 20 6F 66 20 22 2F 62 69 6E 2F 73 68 22 0A 20 20 20 20 78 6F 72 6C 20 25 65 63 78 2C
20 25 65 63 78 20 20 20 20 20 20 20 20 20 23 20 6E 75 6C 6C 20 66 6F 72 20 63 6F 6D 6D 61
6E 64 2D 6C 69 6E 65 20 61 72 67 75 6D 65 6E 74 73 0A 20 20 20 20 78 6F 72 6C 20 25 65
64 78 2C 20 25 65 64 78 20 20 20 20 20 20 20 20 20 23 20 6E 75 6C 6C 20 66 6F 72 20 65 6E
76 69 72 6F 6E 6D 65 6E 74 20 76 61 72 69 61 62 6C 65 73 0A 20 20 20 20 69 6E 74 20 24 30
78 38 30 0A 0A 20 20 20 20 78 6F 72 6C 20 25 65 61 78 2C 20 25 65 61 78 0A 20 20 20 20 69
6E 63 20 25 65 61 78 20 20 20 20 20 20 20 20 20 20 20 20 20 20 20 23 20 73 79 73 63 61
6C 6C 20 6E 75 6D 62 65 72 20 66 6F 72 20 65 78 69 74 0A 20 20 20 20 69 6E 74 20 24 30 78
38 30 0A 0A 2E 73 65 63 74 69 6F 6E 20 2E 64 61 74 61 0A 20 20 20 20 73 68 65 6C 6C 5F 70
61 74 68 3A 20 2E 61 73 63 69 7A 20 22 2F 62 69 6E 2F 73 68 22 0A

Conclusion

In conclusion, the assembly code successfully achieves the goal of invoking the `execve` system call to execute the shell at `/bin/sh`. The shellcode length was determined, and its ASCII representation was examined using the provided Python script. But my shellcode takes a lot of bytes 435 bytes. I will look into more efficient way in the future